

106119029 , OS Lab 6

Dipesh Kafle

Code

```
1 #include <pthread.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include <string.h>
5 #include <sys/types.h>
6 #include <unistd.h>
7 #define min(x, y) x > y ? y : x
8
9 pthread_mutex_t lock;
10
11 struct process_detail {
12     u_int32_t id;
13     u_int32_t arrival;
14     u_int32_t cpu_burst;
15     u_int32_t io_burst;
16     u_int32_t priority;
17     u_int32_t time_quanta;
18     u_int32_t finish_time;
19 };
20
21 struct process_detail *queue[1000];
22
23 struct process_detail_for_io_exec {
24     struct process_detail *process;
25     u_int32_t start_time;
26 };
27
28 struct process_add_timer {
29     struct process_detail *process;
30     u_int32_t *r;
31 };
32
33
34 void *put_in_queue(void *data) {
35     struct process_add_timer *p = (struct process_add_timer *)data;
36     usleep(p->process->arrival * 100);
37     pthread_mutex_lock(&lock);
38     queue[(*(p->r))++] = p->process;
39
40     pthread_mutex_unlock(&lock);
41     printf("Process %d pushed to the end of the queue at %d time \n",
42           p->process->id, p->process->arrival);
43     fflush(stdout);
44 }
45
46 void *sleep_func(void *param) {
47     struct process_detail_for_io_exec p;
48     memcpy(&p, param, sizeof(struct process_detail_for_io_exec));
49     u_int32_t runtime = min(p.process->io_burst, p.process->time_quanta);
50     printf("IO burst started at %d for process %d\n", p.start_time,
51           p.process->id);
52     fflush(stdout);
53     usleep(p.process->io_burst * 100);
54     printf("IO burst finished at %d for process %d\n", p.start_time + runtime,
55           p.process->id);
56     fflush(stdout);
57     p.process->io_burst -= runtime;
58     p.process->time_quanta -= 5;
59     if (p.process->cpu_burst == 0 && p.process->io_burst == 0) {
60         printf("Process %d has completed IO and CPU burst and is no more active\n",
61               p.process->id);
62         p.process->finish_time = p.start_time + runtime;
63     }
64     return NULL;
65 }
```

```

26 int main() {
25     pthread_t threads[100];
24     int thread_count = 0;
23
22     struct process_detail p1 = {1, 0, 120, 10, 5, 50, 0};
21     struct process_detail p2 = {2, 0, 55, 45, 5, 50, 0};
20     struct process_detail p3 = {3, 0, 150, 10, 5, 50, 0};
19     struct process_detail p4 = {4, 0, 65, 35, 5, 50, 0};
18     struct process_detail p5 = {5, 320, 25, 30, 5, 50, 0};
17
16     queue[0] = &p1;
15     queue[1] = &p2;
14     queue[2] = &p3;
13     queue[3] = &p4;
12     u_int32_t l = 0;
11     u_int32_t r = 4;
10     u_int32_t seconds = 0;
9     u_int32_t last_stop = 0;
8
7     struct process_add_timer p = {&p5, &r};
6     // create a thread which will put process 5 to the queue
5     pthread_create(threads + (thread_count++), NULL, put_in_queue, &p);
4
3     int io = 0;
2     while (r != l) {
1
93     u_int32_t runtime = min(queue[l]→cpu_burst, queue[l]→time_quanta);
1     seconds += runtime;
2     usleep(runtime * 100);
3     queue[l]→time_quanta =
4         min(queue[l]→time_quanta + (runtime == queue[l]→time_quanta ? 10 : 0),
5             100);
6
100    queue[l]→cpu_burst -= runtime;
1     queue[l]→priority++;
2     printf("Process %d executed for %d time(starting: %d, end:%d)\n ",
3         queue[l]→id, runtime, seconds - runtime, seconds);
4     fflush(stdout);
5
6     if (queue[l]→io_burst) {
7         struct process_detail_for_io_exec data = {queue[l], seconds};
8         pthread_create(threads + thread_count, NULL, sleep_func, &data);
9         thread_count++;
10    }
11
12    if (queue[l]→cpu_burst) {
13        printf("    Process %d pushed back to the queue\n", queue[l]→id);
14        queue[r++] = queue[l++];
15    } else {
16        printf("    CPU Execution time for process %d is over\n", queue[l]→id);
17        // if io_burst and cpu_burst both completed process is done
18        if (queue[l]→io_burst == 0) {
19            printf(
20                "    Process %d has completed IO and CPU burst and is no more in "
21                "queue\n",
22                queue[l]→id);
23            queue[l]→finish_time = seconds;
24        }
25        l++;
26    }
27    fflush(stdout);
28 }
29
130 printf("\n\n\n");
1     struct process_detail *all_processes[5] = {&p1, &p2, &p3, &p4, &p5};
2
3     double turnaroundtime = 0;
4     for (int i = 0; i < 5; i++) {
5         printf("Process %d arrived at %d and finished at %d\n",
6             all_processes[i]→id, all_processes[i]→arrival,
7             all_processes[i]→finish_time);
8         turnaroundtime +=
9             (all_processes[i]→finish_time - all_processes[i]→arrival);
10    }
11    printf("\nThe average turnaroundtime is %f\n", turnaroundtime / 5);
12 }

```


Output

```
Process 1 executed for 50 time(starting: 0, end:50)
  Process 1 pushed back to the queue
IO burst started at 50 for process 1
IO burst finished at 60 for process 1
Process 2 executed for 50 time(starting: 50, end:100)
  Process 2 pushed back to the queue
IO burst started at 100 for process 2
IO burst finished at 145 for process 2
Process 3 executed for 50 time(starting: 100, end:150)
  Process 3 pushed back to the queue
IO burst started at 150 for process 3
IO burst finished at 160 for process 3
Process 4 executed for 50 time(starting: 150, end:200)
  Process 4 pushed back to the queue
IO burst started at 200 for process 4
IO burst finished at 235 for process 4
Process 1 executed for 55 time(starting: 200, end:255)
  Process 1 pushed back to the queue
Process 2 executed for 5 time(starting: 255, end:260)
  CPU Execution time for process 2 is over
  Process 2 has completed IO and CPU burst and is no more in queue
Process 5 pushed to the end of the queue at 320 time
Process 3 executed for 55 time(starting: 260, end:315)
  Process 3 pushed back to the queue
Process 4 executed for 15 time(starting: 315, end:330)
  CPU Execution time for process 4 is over
  Process 4 has completed IO and CPU burst and is no more in queue
Process 1 executed for 15 time(starting: 330, end:345)
  CPU Execution time for process 1 is over
  Process 1 has completed IO and CPU burst and is no more in queue
Process 5 executed for 25 time(starting: 345, end:370)
  CPU Execution time for process 5 is over
IO burst started at 370 for process 5
IO burst finished at 400 for process 5
Process 5 has completed IO and CPU burst and is no more active
Process 3 executed for 45 time(starting: 370, end:415)
  CPU Execution time for process 3 is over
  Process 3 has completed IO and CPU burst and is no more in queue

Process 1 arrived at 0 and finished at 345
Process 2 arrived at 0 and finished at 260
Process 3 arrived at 0 and finished at 415
Process 4 arrived at 0 and finished at 330
Process 5 arrived at 320 and finished at 400

The average turnaroundtime is 286.000000
```

Gantt Chart



- I have a queue for CPU and I've simulated IO using threads. After one quanta is over for a process, it will wait for IO in a separate thread for the `io_burst` amount of time. That way I have simulated IO and CPU times.
- `p1, p2, p3` and `p4` are put into the queue at 0 time.
- At 0, `P1` starts the execution and continues till 50 time units. Then it will stop there and add 10 time units to its quanta because it executed for a whole quanta without waiting for IO. It will also spawn a thread for IO waiting, which will sleep for 10 time units and terminate after modifying the values and printing. At `t=50`, `p1.time_quanta=60`.
- At 50-60 time interval in the IO thread, `P1` has finished its IO and printed. Since IO was done, time quanta of `P1` is reduced. At `t=60`, `p1.time_quanta = 60-5=55`.
- At `t=50`, `P2` starts its CPU execution and continues until 100 in the CPU ready queue. It then spawns a thread to do IO (which goes on for 45 seconds from 100 to 145 time units). So for the next CPU cycle of `P2`, it will have **55 time units** as it did one whole quanta of CPU bound work and also did IO work. So `50+10-5`.
- At `t=100`, `P3` starts executing till 150 and spawns a thread for IO which will go on for 10 time units (till 160). It's next time quanta will also be **55 time units**.
- At `t=150`, `P4` starts off and goes on for a quanta till 200 and then spawns a thread which goes on for 35 time units. It's time quanta for next cycle will also be **55 time units**.
- At `t=200`, `P1` is scheduled and it starts off and works for 55 time units (one full quanta) and since IO is over for `P1`, it won't spawn a new thread. It's time quanta is increased to **65 time units**.

- At t=255 , P2 starts off and goes on till 260 and finishes all its work and prints that. Hands the control to scheduler at 260.
- At t=260, P3 starts and goes on for 55 time units(its time quanta) and is pushed back to the queue again. No IO is done as its IO time is over already.
- At t=315, P4 starts off and goes on for 15 time units until it completes all its work. Wont be available for scheduling again.
- At t=330, P1 is given the CPU and it will complete all its work at 345 and hand the control to the scheduler.
- At t=345, P5 starts the execution for the first time ever and finishes CPU work by 370. Spawns a new thread at t=370 and hands the control to scheduler. The IO will complete at t=400, after which process P5 will have finished all its work.
- At t=370, CPU is handed over to P3 which will complete everything by 415 and now all our processes have been completed.
- From the output we can see that the Average Turnaround time is 286 time units.

Remarks

- It is obvious by how the priority is changing this scheudling method will favor a more CPU bounded process.

Alternative approach for the sake of comparision

- I have also made another scheduler, A classic round robin scheduler. In that, CPU time and IO time are treated as same for the sake of simplicity. When I schedule these processes with that scheduler, the processes are finishing at **545 time units** which is much higher. Obviously the turnaround time is more for it (387 time units).
- Output for normal RR scheduler

```

Process 1 executed for one time quanta of 50 time unit(without IO), pushed back to the end of queue at 50
Process 2 executed for one time quanta of 50 time unit(without IO), pushed back to the end of queue at 100
Process 3 executed for one time quanta of 50 time unit(without IO), pushed back to the end of queue at 150
Process 4 executed for one time quanta of 50 time unit(without IO), pushed back to the end of queue at 200
Process 1 executed for one time quanta of 60 time unit(without IO), pushed back to the end of queue at 260
Process 2 finished at 310 time unit with IO
Process 5 pushed to the end of queue at 320
Process 3 executed for one time quanta of 60 time unit(without IO), pushed back to the end of queue at 370
Process 4 finished at 420 time unit with IO
Process 1 finished at 440 time unit with IO
Process 5 executed for one time quanta of 50 time unit(with IO), pushed back to the end of queue at 490
Process 3 finished at 540 time unit with IO
Process 5 finished at 545 time unit with IO

Process 1 arrived at 0 and finished at 440
Process 2 arrived at 0 and finished at 310
Process 3 arrived at 0 and finished at 540
Process 4 arrived at 0 and finished at 420
Process 5 arrived at 320 and finished at 545

The average turnaroundtime is 387.000000

```

- Code for normal RR scheduler

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <sys/types.h>
4 #include <unistd.h>
5
6 struct process_detail
7 {
8     u_int32_t id;
9     u_int32_t arrival;
10    u_int32_t cpu_burst;
11    u_int32_t io_burst;
12    u_int32_t priority;
13    u_int32_t time_quanta;
14    u_int32_t finish_time;
15 };
16
17 int main() {
18     struct process_detail p1 = {1, 0, 120, 10, 5, 50, 0};
19     struct process_detail p2 = {2, 0, 55, 45, 5, 50, 0};
20     struct process_detail p3 = {3, 0, 150, 10, 5, 50, 0};
21     struct process_detail p4 = {4, 0, 65, 35, 5, 50, 0};
22     struct process_detail p5 = {5, 320, 25, 30, 5, 50, 0};
23
24     struct process_detail *queue[1000] = {NULL};
25     queue[0] = &p1;
26     queue[1] = &p2;
27     queue[2] = &p3;
28     queue[3] = &p4;
29     int l = 0;
30     int r = 4;
31     /* u_int32_t time_quanta = 50; */
32     u_int32_t seconds = 0;
33     u_int32_t last_stop = 0;
34 }

```

```

34 int io = 0;
1 while (r != l) {
2
3     if (queue[l]→cpu_burst) {
4         // we'll finish the cpu burst first
5         queue[l]→cpu_burst--;
6     } else {
7         // then we move on to io_burst
8         if (queue[l]→io_burst) {
9             queue[l]→io_burst--;
10            io = 1;
11        } else {
12            // if there's no io_burst as well, we don't put the process in queue
13            // anymore as it is finished
14            queue[l]→finish_time = seconds;
15            last_stop = seconds;
16            printf("Process %d finished at %d time unit with%s\n", queue[l]→id,
17                queue[l]→finish_time, (io == 0) ? "out IO" : " IO");
18            io = 0;
19            l++;
20            continue;
21        }
22    }
23
24    seconds++;
25
60 if (seconds == 320) {
1     queue[r++] = &p5;
2     printf("Process %d pushed to the end of queue at %d\n", queue[r - 1]→id,
3         queue[r - 1]→arrival);
4 }
5 // if the time quanta is finished
6 if (seconds - last_stop == queue[l]→time_quanta) {
7     printf("Process %d executed for one time quanta of %d time unit(with%s), "
8         "pushed "
9         "back to the end of queue at %d\n",
10        queue[l]→id, queue[l]→time_quanta, (io == 0) ? "out IO" : " IO",
11        seconds);
12    if (io) {
13        queue[l]→time_quanta -= 5;
14    } else {
15        queue[l]→time_quanta += (queue[l]→time_quanta == 100 ? 0 : 10);
16    }
17    queue[r++] = queue[l++];
18    last_stop = seconds;
19 }
20 }
21 printf("\n\n\n");
22 struct process_detail *all_processes[5] = {&p1, &p2, &p3, &p4, &p5};
23
24 double turnaroundtime = 0;
25 for (int i = 0; i < 5; i++) {
26     printf("Process %d arrived at %d and finished at %d\n",
27         all_processes[i]→id, all_processes[i]→arrival,
28         all_processes[i]→finish_time);
29     turnaroundtime +=
30         (all_processes[i]→finish_time - all_processes[i]→arrival);
31 }
32 printf("\nThe average turnaroundtime is %f\n", turnaroundtime / 5);
33 }

```